Chapter 10: Factorial ANOVA

Statistics for Environmental Science

Mallory Barnes

Rendered October 20, 2026

This PDF was rendered on October 20, 2026. The book is updated during the semester; the latest version of this chapter is at https://malloryb.github.io/statistics_E538/10-FactorialANOVA.html.

Statistics for Environmental Science · CC BY-SA 4.0 · Credits and license: https://malloryb.github.io/statistics_E538/

Table of contents

Factorial basics **2**
Factorial notation . . . 3

Pollution, flow, and stream life **3**
Cell means and marginal means . . . 4

Why factorial designs matter **5**
Main effects and interactions . . . 5

What a 2x2 actually tests **6**
Hypotheses . . . 7

What the ANOVA compares **7**

Computing a factorial ANOVA by hand **8**
SS Total . . . 8
SS Pollution . . . 8
SS Flow . . . 9
SS Interaction . . . 9
SS Error . . . 10
Degrees of freedom, mean squares, and F . . . 10
Check your work . . . 10
Checking the assumptions . . . 11
Effect size in a factorial design . . . 11

Interpreting main effects and interactions **12**
Reading patterns in a 2x2 . . . 12
Simple effects . . . 15
Order of interpretation . . . 16
A consistent main effect and an interaction . . . 16

An "averaged" main effect driven by an interaction . . . 16
A main effect masked by an interaction . . . 18
Reporting the stream result . . . 19

More complicated designs **20**
3x2 design . . . 21

Designs beyond this course **23**
2x2x2 (three-factor) designs . . . 23
Mixed designs . . . 23

Chapter Summary **24**

In science, outcomes rarely depend on just one thing. Climate, soil type, and water availability all act together to influence ecosystems in ways a single-factor design can't capture. *Factorial designs* let us study several factors, and how they interact with each other, at once.

Chapter 8's one-way ANOVA tested a single factor. This chapter covers how to compute a factorial ANOVA, including a full by-hand walkthrough of the sums-of-squares arithmetic, and how to interpret the result. With more than one factor, interpretation is harder, because an *interaction* changes what every other number in the results table means.

Factorial basics

Imagine investigating the growth rate of a plant species, with independent variables like soil pH or water variability. Each added IV multiplies the number of conditions you can observe, and with it, the questions you can ask about how those factors work together. Table 1 shows three designs built around the same plant-growth question.

Table 1: Each added IV or level changes the test.

Design	Example question	Test
1 IV, 2 levels (soil pH: acidic, neutral)	How does soil pH affect plant growth?	t-test
1 IV, 3 levels (water: low, medium, high)	How does water availability affect plant growth?	One-way ANOVA
2 IVs, 2 levels each (soil pH $\times$ water: low, high)	How do soil pH and water availability together affect plant growth?	Factorial ANOVA

Now we have two IVs, soil pH and water availability, each with two levels. A **factorial** design has two or more IVs, and each of its conditions is a combination of one level from each IV.

A **fully factorial** design includes every combination of the levels of its IVs, so each level of one IV occurs with each level of every other IV. A design with this property is also called fully crossed. A factorial design that leaves some combinations of levels unmeasured is only partly crossed. Every design in this chapter is fully factorial, so the rest of the chapter uses factorial to mean fully factorial.

Factorial notation

Each combination of soil type and water availability creates a unique environment for the plants, giving four distinct scenarios to measure growth rates: 2 IVs, each with 2 levels, for 2x2 = 4 conditions total. The notation directly tells you the number of conditions.

The notation system indicates the structure of the factorial design. Each IV gets a number representing its levels. Here are a few examples:

- 2x2: Two IVs, each with two levels (yields four unique conditions).
- 2x3: Two IVs, with the first having two levels and the second three (yields six unique conditions).
- 3x2: Similar to the 2x3, but the first IV has three levels, and the second has two, also giving us six unique conditions.
- 4x4: Two IVs, each with four levels, resulting in sixteen conditions.

Pollution, flow, and stream life

The rest of this chapter uses one example throughout.

The question. Streams are affected by pollution, but also by how fast the water moves. Fast-moving water flushes contaminants downstream and re-oxygenates the channel. Slow water lets them sit. So pollution might do more damage in a sluggish stream than a fast one. To measure that damage, we need to operationalize stream health.

The dependent variable is macroinvertebrate species count, the number of distinct kinds of aquatic insects, snails, worms, and crustaceans that live in a stream. This is a standard measure of stream health. These animals are in the middle of the food web, many are pollution-sensitive, and they can't flee a degraded stream the way fish can. More species present means a healthier stream.

The design. We run an experiment in artificial stream channels, each assigned one pollution condition and one flow rate. The species counts in this chapter are invented, chosen so the arithmetic stays easy to follow. Figure 1 summarizes it.

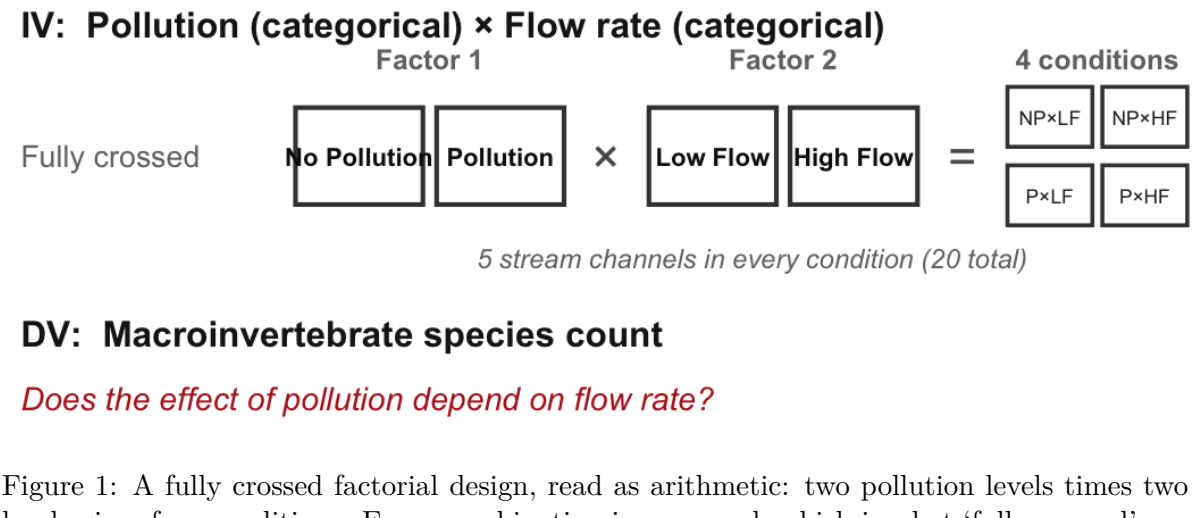


Figure 1: A fully crossed factorial design, read as arithmetic: two pollution levels times two flow levels gives four conditions. Every combination is measured, which is what 'fully crossed' means.

With two IVs of two levels each, this is a 2x2 design. Reading the schematic left to right also shows why the notation works the way it does. Crossing two levels with two levels *is* the multiplication that gives four conditions, and "fully crossed" means no cell is left empty.

That four-cell grid resembles the chi-square table from Chapter 7, but the cells mean different things. There, the cells held *counts* of city blocks that already had a given land use and either did or did not top 40°C, and nobody put them there. Here, the cells are *conditions we assign channels to*, and within each one we then measure something else entirely, the species count.

Predictions. No pollution plus high flow should support the most species, and pollution plus low flow the fewest. The other two combinations are the informative ones. They show whether pollution matters only when flow is low, with fast water washing its effect away, or hurts a stream no matter how fast the water moves. A one-factor study could not tell these apart, and the interaction test in a factorial experiment can.

Table 2 lists the species count in every channel.

Table 2: Macroinvertebrate species count in each of the 20 stream channels, five per condition.

Channel	No Pollution, Low Flow	Pollution, Low Flow	No Pollution, High Flow	Pollution, High Flow
1	10	5	12	9
2	8	4	13	8
3	11	3	14	10
4	9	4	11	11
5	10	2	13	12

Cell means and marginal means

Each condition, such as no pollution with low flow, is a cell, the same word Chapter 7 used for one combination in a table of counts. A **cell mean** is the mean of the scores in one cell. Here, each cell mean averages the species counts of the five stream channels assigned to that condition. A **marginal mean** is the mean for one level of a factor, averaged across the levels of the other factor. For example, the no-pollution marginal mean averages the two no-pollution cells, one at each flow rate. Table 3 shows the means from the experiment this chapter analyzes.

Table 3: Mean macroinvertebrate species count. The four inner numbers are cell means; the last row and column hold the marginal means, and the corner is the grand mean.

	Low Flow	High Flow	Marginal mean
No Pollution	9.6	12.6	11.1
Pollution	3.6	10.0	6.8
Marginal mean	6.6	11.3	8.95

The pollution marginal means (11.1 and 6.8) average each row across both flow rates. The flow marginal means (6.6 and 11.3) average each column across both pollution conditions. The corner, 8.95, is the grand mean of all 20 channels.

Why factorial designs matter

Factorial designs let us ask more nuanced questions than a single-factor study can. Suppose we ignored flow entirely and just compared polluted to unpolluted channels. That compares the two pollution marginal means, 11.1 vs. 6.8 species, shown in Figure 2.

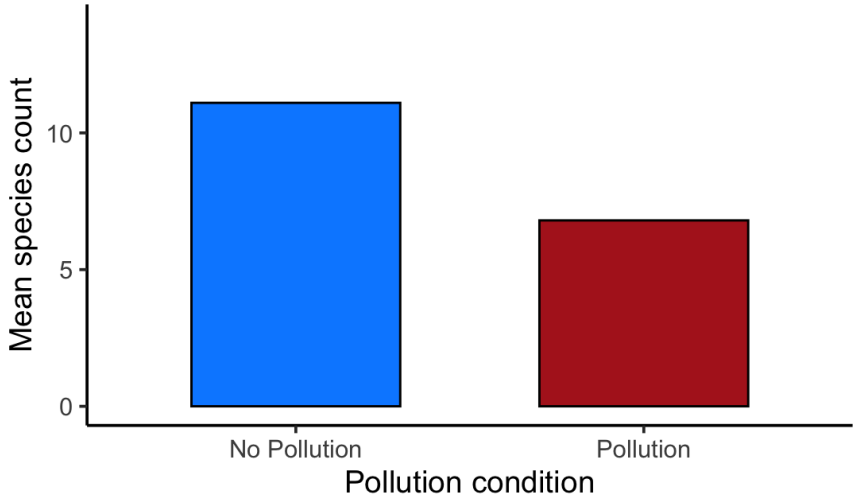


Figure 2: Mean macroinvertebrate species count under no pollution vs. pollution, averaged across flow rate (the pollution marginal means).

A factorial design lets us ask whether the effect of pollution itself depends on flow, which a one-factor design can't. With flow rate added as a second IV, there are four cell means instead of two, so we can compare polluted and unpolluted channels *within* each flow rate, shown in Figure 3.

Under low flow, species count drops sharply with pollution (9.6 to 3.6, a difference of 6.0). Under high flow, the drop is smaller (12.6 to 10.0, a difference of 2.6). The effect of pollution depends on flow rate.

Main effects and interactions

A factorial design gives you two kinds of results: main effects and interactions. How many of each depends on the number of IVs.

A **main effect** is the difference between the marginal means of one IV. Each IV has one main effect, so a 2x2 design has two, one for Pollution and one for Flow. A 2x2x2 design has three main effects, one per IV. A 3x3x3 design also has three, since the number of main effects tracks the number of IVs, not their levels.

A **simple effect** is the effect of one IV at a single level of the other IV, such as the pollution effect under low flow. Figure 3 shows two of them, one per flow rate, and the main effect of pollution is their average: $(6.0 + 2.6)/2 = 4.3$, the same 4.3 you get by subtracting the pollution marginal means.

An **interaction** occurs when the effect of one IV depends on the level of another. Figure 3 shows one, since the simple effect of pollution changed with flow.

Equivalently, an interaction is a **difference between differences**, which is what the arithmetic computes. Each simple effect of pollution is a difference (unpolluted minus polluted). If that difference itself changes across flow rates, that's an interaction.

Table 4 collects the four quantities a 2x2 design gives, with the stream experiment's value for each.

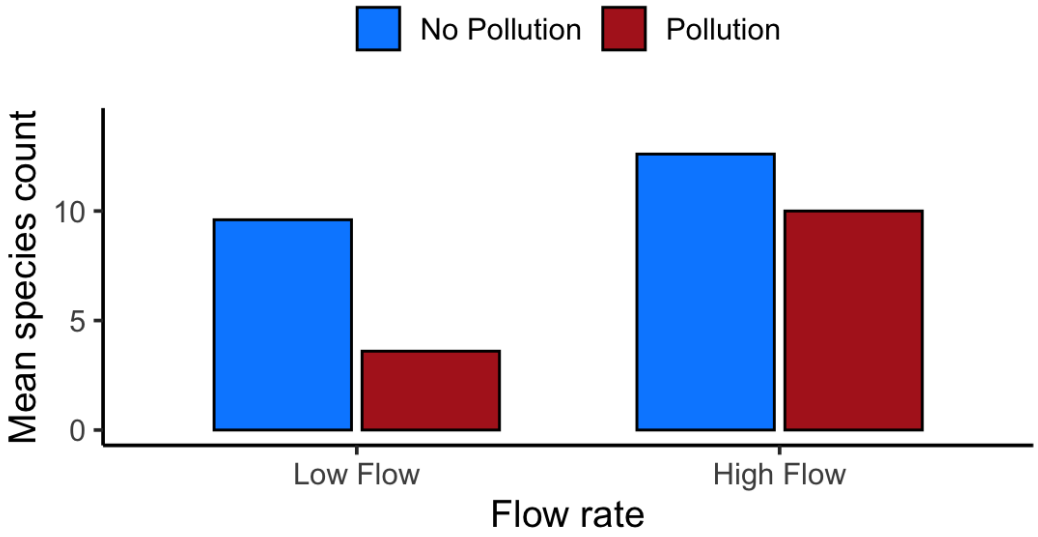


Figure 3: The effect of pollution within each flow rate (the gap between the two bars) is larger under low flow than under high flow.

Table 4: Cell mean, main effect, simple effect and interaction in the stream experiment.

Term	Definition	Stream experiment
Cell mean	The mean of the channels in one condition	No pollution with low flow: 9.6 species
Main effect	The difference between one factor's marginal means, averaged across the levels of the other factor	Pollution: $11.1 - 6.8 = 4.3$ species
Simple effect	The effect of one factor at a single level of the other	Pollution under low flow: 6.0 species. Under high flow: 2.6 species.
Interaction	The effect of one factor depends on the level of the other. The simple effects differ, a difference between differences.	$6.0 - 2.6 = 3.4$ species

The number of possible interactions scales with the number of IVs. A 2x2 design has exactly one interaction, the combined effect of the two IVs. A design with three IVs (A, B, C) has three two-way interactions (AB, AC, BC) and one three-way interaction (ABC).

What a 2x2 actually tests

A two-way ANOVA tests three effects and therefore has three null hypotheses:

1. a main effect of pollution,
2. a main effect of flow, and
3. a pollution $\times$ flow interaction.

The hypotheses are about the population mean species count for each condition, written μ. The cell means computed from the 20 channels, written $\bar{Y}$ later in this chapter, estimate them. We can represent the four conditions and their mean species counts as:

Pollution	Flow	Mean species count
No Pollution	Low Flow	$\mu_{\text{NoPoll,Low}}$
No Pollution	High Flow	$\mu_{\text{NoPoll,High}}$
Pollution	Low Flow	$\mu_{\text{Poll,Low}}$
Pollution	High Flow	$\mu_{\text{Poll,High}}$

Hypotheses

Main effect of pollution

- **Null (H_0):** There is no effect of pollution on mean species count.
- **Alternative (H_a):** Pollution affects mean species count.

Main effect of flow

- **Null (H_0):** There is no effect of flow rate on mean species count.
- **Alternative (H_a):** Flow rate affects mean species count.

Pollution × flow interaction

- **Null (H_0):** The effect of pollution is the same at both flow rates. *(In symbols:* $(\mu_{\text{NoPoll,Low}} - \mu_{\text{Poll,Low}}) = (\mu_{\text{NoPoll,High}} - \mu_{\text{Poll,High}})$.)
- **Alternative (H_a):** The effect of pollution differs across flow rates.

Under the interaction null, the gap between polluted and unpolluted channels is the same whether the water is moving fast or slow. If that gap changes, for example if pollution does far more damage under low flow, you've found an interaction.

> **Note:** The interaction null for any two factors
>
> You'll sometimes see this written in abstract terms, with "Factor A" and "Factor B." For example: $H_0 : (\mu_{A1,B1} - \mu_{A2,B1}) = (\mu_{A1,B2} - \mu_{A2,B2})$. The logic is the same, with A and B standing for your own variables.

What the ANOVA compares

An interaction is a difference between differences. Table 3 gives the four cell means to compute it with, before any sums of squares. Writing a simple effect as Δ, with the level it is computed at after the bar:

$$\Delta_{\text{Pollution|Low}} = \bar{Y}_{\text{NoPoll,Low}} - \bar{Y}_{\text{Poll,Low}} = 9.6 - 3.6 = 6.0$$

$$\Delta_{\text{Pollution|High}} = \bar{Y}_{\text{NoPoll,High}} - \bar{Y}_{\text{Poll,High}} = 12.6 - 10.0 = 2.6$$

Those two numbers aren't the same, so the effect of pollution isn't constant across flow rates, which is exactly what an interaction means. The gap between them, $6.0 - 2.6 = 3.4$, is a first, informal look at the interaction itself. The formal test below tells us whether a gap that size is larger than chance would produce.

Computing a factorial ANOVA by hand

In our between-subjects 2x2 design, different stream channels appear in each condition: 5 stream channels assigned to each condition, 20 channels total, no repeats. The logic is the same as the one-way ANOVA in Chapter 8, which partitions the total variance into parts. For a 2x2 design, there's now one more part, the interaction:

$$SS_{\text{Total}} = SS_{\text{Pollution}} + SS_{\text{Flow}} + SS_{\text{Interaction}} + SS_{\text{Error}}$$

The tables below build up each piece the same way the one-way ANOVA chapter did: one row per channel, showing what that channel's score gets replaced with, its distance from the grand mean, and that distance squared.

SS Total

SS_{Total} captures all of the variation in the data. It squares the difference between each channel's score and the grand mean and adds them up. The table shows the first few channels and the last, since the pattern is the same all the way down.

Table 5: Deviations from the grand mean

Pollution	Flow	Species	Deviation from grand mean	Squared deviation
No Pollution	Low Flow	10	1.05	1.1
No Pollution	Low Flow	8	-0.95	0.9
No Pollution	Low Flow	11	2.05	4.2
…	…	…	…	…
Pollution	High Flow	12	3.05	9.3
Sums		179	0	242.95
Means		8.95		

The mean of all 20 scores is the grand mean, 8.95. The deviation column holds every channel's distance from it, and the last column squares those distances so they don't cancel (the same reason we squared deviations for the one-way ANOVA). Summing the squared column gives $SS_{\text{Total}} = 242.95$.

SS Pollution

$SS_{\text{Pollution}}$ isolates the variation we can explain by pollution condition alone. To compute it, replace every score with its pollution marginal mean, then find that mean's distance from the grand mean.

Table 6: Pollution marginal means replace the scores

Pollution	Flow	Species	Pollution marginal mean	Marginal mean minus grand mean	Squared
No Pollution	Low Flow	10	11.1	2.15	4.62
No Pollution	Low Flow	8	11.1	2.15	4.62
No Pollution	Low Flow	11	11.1	2.15	4.62
…	…	…	…	…	…
Pollution	High Flow	12	6.8	-2.15	4.62

Pollution	Flow	Species	Pollution marginal mean	Marginal mean minus grand mean	Squared
Sums		179		0	92.45

Every no-pollution channel is replaced by the no-pollution marginal mean, 11.1, and every pollution channel by the pollution marginal mean, 6.8. Squaring and summing those means' distances from the grand mean gives $SS_{\text{Pollution}} = 92.45$, the variation explained by pollution condition alone.

SS Flow

SS_{Flow} works the same way, grouping by flow rate instead of pollution condition.

Table 7: Flow marginal means replace the scores

Pollution	Flow	Species	Flow marginal mean	Marginal mean minus grand mean	Squared
No Pollution	Low Flow	10	6.6	-2.35	5.52
No Pollution	Low Flow	8	6.6	-2.35	5.52
No Pollution	Low Flow	11	6.6	-2.35	5.52
…	…	…	…	…	…
Pollution	High Flow	12	11.3	2.35	5.52
Sums		179		0	110.45

Every low-flow channel is replaced by the low-flow marginal mean, 6.6, and every high-flow channel by the high-flow marginal mean, 11.3. This gives $SS_{\text{Flow}} = 110.45$.

SS Interaction

The interaction is the one new piece a factorial ANOVA adds. It asks whether the four cell means behave differently than the two main effects alone would predict.

If the main effects told the whole story, each cell mean would be the grand mean, shifted by that cell's pollution effect and its flow effect. Take the no-pollution, low-flow cell. The no-pollution marginal mean is $11.1 - 8.95 = 2.15$ above the grand mean, and the low-flow marginal mean is $8.95 - 6.6 = 2.35$ below it, so the two main effects predict $8.95 + 2.15 - 2.35 = 8.75$. The observed cell mean is 9.6. The leftover, 0.85, is the part of that cell the two main effects can't explain. Those leftovers, one per cell, are what the interaction captures.

To compute it, replace each score with its own cell mean, then subtract that channel's pollution marginal mean and flow marginal mean, and add back the grand mean once (to avoid subtracting it twice):

$$\bar{Y}_{\text{cell}} - \bar{Y}_{\text{Pollution marginal}} - \bar{Y}_{\text{Flow marginal}} + \bar{Y}_{\text{grand}}$$

Table 8: What is left of each cell mean after both main effects

Pollution	Flow	Species	Cell mean	Interaction leftover	Squared
No Pollution	Low Flow	10	9.6	0.85	0.72
No Pollution	Low Flow	8	9.6	0.85	0.72

Pollution	Flow	Species	Cell mean	Interaction leftover	Squared
No Pollution	Low Flow	11	9.6	0.85	0.72
…	…	…	…	…	…
Pollution	High Flow	12	10	0.85	0.72
Sums		179		0	14.45

Squaring and summing gives $SS_{\text{Interaction}} = 14.45$.

SS Error

The last piece follows by subtraction, since every other term in the identity is already known:

$$SS_{\text{Error}} = SS_{\text{Total}} - SS_{\text{Pollution}} - SS_{\text{Flow}} - SS_{\text{Interaction}}$$

$$SS_{\text{Error}} = 242.95 - 92.45 - 110.45 - 14.45 = 25.6$$

SS_{Error} is the same quantity Chapter 8 called SS_{Within} in the one-way ANOVA, the spread of channels around their own cell mean. It is the variation none of the three effects can explain.

Degrees of freedom, mean squares, and *F*

Each sum of squares becomes a mean square the same way it did in Chapter 8, by dividing by its degrees of freedom. With a levels of pollution, b levels of flow, and N channels in all:

- $df_{\text{Pollution}} = a - 1 = 1$
- $df_{\text{Flow}} = b - 1 = 1$
- $df_{\text{Interaction}} = (a-1)(b-1) = 1$
- $df_{\text{Error}} = N - ab = 20 - 4 = 16$, one lost for each of the four cell means
- $df_{\text{Total}} = N - 1 = 19$, and $1 + 1 + 1 + 16 = 19$

Each effect gets its own F, and all three share the same denominator:

$$F_{\text{Pollution}} = \frac{MS_{\text{Pollution}}}{MS_{\text{Error}}}, \qquad F_{\text{Flow}} = \frac{MS_{\text{Flow}}}{MS_{\text{Error}}}, \qquad F_{\text{Interaction}} = \frac{MS_{\text{Interaction}}}{MS_{\text{Error}}}$$

Here $MS_{\text{Error}} = 25.60/16 = 1.60$. Every effect has 1 degree of freedom, so each effect's MS equals its SS, and for pollution $F = 92.45/1.60 = 57.78$.

Check your work

R's `aov()` builds the same table directly. The `*` in the formula tells R to include both main effects and their interaction. R writes the interaction row as `Pollution:Flow` and the error row as `Residuals`.

```r
all_df <- data.frame(
  Number_species = c(10, 8, 11, 9, 10,  5, 4, 3, 4, 2,
                     12, 13, 14, 11, 13,  9, 8, 10, 11, 12),
  Pollution = rep(c("No Pollution", "Pollution", "No Pollution", "Pollution"), each = 5),
  Flow      = rep(c("Low Flow", "Low Flow", "High Flow", "High Flow"), each = 5)
)
aov_fit <- aov(Number_species ~ Pollution * Flow, data = all_df)
summary(aov_fit)
```

Run these lines in R to reproduce the table below.

Source	df	SS	MS	F	p
Pollution	1	92.45	92.45	57.78	< 0.001
Flow	1	110.45	110.45	69.03	< 0.001
Pollution:Flow	1	14.45	14.45	9.03	0.008
Residuals	16	25.60	1.60		

The SS column matches the hand-calculated values above: 92.45 for Pollution, 110.45 for Flow, 14.45 for the interaction, 25.6 for Error. All three effects are significant here, including the interaction we previewed informally. The difference in the pollution effect between the two flow rates is larger than chance would produce.

This example is balanced (Chapter 8), with five stream channels in every condition. With unequal numbers, two things change. The sums of squares depend on the order of terms in `aov()`, because R computes them sequentially, and the marginal means become model-based estimates (estimated marginal means) rather than simple averages of the scores. Balanced designs avoid both complications, so plan equal sample sizes when you can.

Checking the assumptions

A factorial ANOVA has the same test-specific assumptions as the one-way ANOVA in Chapter 8: similar spread in every cell and approximately normal residuals. Check them the same way, with Levene's test across the four cells and a Q-Q plot or Shapiro-Wilk test on the residuals of `aov_fit`.

Effect size in a factorial design

The one-way ANOVA in Chapter 8 had one effect and so needed one effect size, $\eta^2 = SS_{\text{Between}}/SS_{\text{Total}}$. A 2x2 runs three tests, so it needs three effect sizes, and the sums of squares we just computed by hand are all the ingredients required.

The direct extension keeps SS_{Total} in the denominator for every effect:

$$\eta^2 = \frac{SS_{\text{Effect}}}{SS_{\text{Total}}}$$

Each effect then reports its share of *all* the variation in the study, so the three values plus error's share add to 1. That is easy to interpret, but adding a second factor to a study inflates SS_{Total}, which shrinks the η^2 of a factor whose behavior did not change at all. The same effect looks smaller purely because you studied something else alongside it.

Partial eta-squared (η_p^2) is the share of the variation left after setting the other effects aside that one effect explains. Its denominator is that effect's SS plus SS_{Error}:

$$\eta_p^2 = \frac{SS_{\text{Effect}}}{SS_{\text{Effect}} + SS_{\text{Error}}}$$

Because each effect gets its own denominator, partial values do not sum to 1 and are always at least as large as the corresponding η^2.

The distinction never came up in Chapter 8, because a one-way ANOVA has only one effect, so there are no other factors to set aside, and $SS_{\text{Effect}} + SS_{\text{Error}}$ *is* SS_{Total}. The two formulas give the same number there. They only diverge once a design has more than one factor.

Table 9: Eta-squared and partial eta-squared for each effect in the 2x2 design.

Effect	SS	Eta-squared	Partial eta-squared
Pollution	92.45	0.381	0.783
Flow	110.45	0.455	0.812
Pollution x Flow	14.45	0.059	0.361

Flow explains 45% of the total variation, pollution 38%, and the interaction 6%. The interaction is significant, and it is also the smallest of the three effects, a common combination, because significance and magnitude answer different questions. The interaction explains little of the total variation, but it changes what the pollution main effect means.

Say which one you report. Most statistical software, including SPSS and R's `effectsize` package, gives you partial η_p^2 by default, and it is the more common choice in published work. Either is defensible, but say which you used, because the two can differ substantially: pollution is $\eta^2 = 0.38$ but $\eta_p^2 = 0.78$ here. Reporting the larger number without labeling it overstates the effect. Never compare a partial value from one study against an unpartialed value from another.

Interpreting main effects and interactions

With the ANOVA table computed, the next step is to interpret it.

It's convenient to think of a main effect as one manipulation's consistent influence. An interaction means the effect of one IV depends on the level of another. So when there is an interaction, at least one main effect is not consistent across conditions.

Reading patterns in a 2x2

In any 2x2 ANOVA, there are eight general outcome patterns, different combinations of main effects and interactions. You don't need to memorize all eight, but you should learn to read them in a line plot:

Pattern	Main Effect: Pollution	Main Effect: Flow	Interaction
1			
2	X		
3	X		X
4	X	X	
5	X	X	X

Pattern	Main Effect: Pollution	Main Effect: Flow	Interaction
6		X	
7		X	X
8			X

Figure 4 shows the eight patterns as bar graphs, which show differences in means directly, and Figure 5 below shows the same patterns as line graphs, which make interactions easier to spot. These are abstract illustrations of every combination that can occur, so a few of them describe patterns you would not expect ecologically. Use them to learn to recognize the shapes.

Every figure in this chapter puts flow rate on the x-axis and draws one line (or bar) per pollution condition, blue for no pollution and red for pollution.

The eight panels below are numbered to match the table above, so panel "Pattern 3" in the figure is row 3 in the table, and so on.

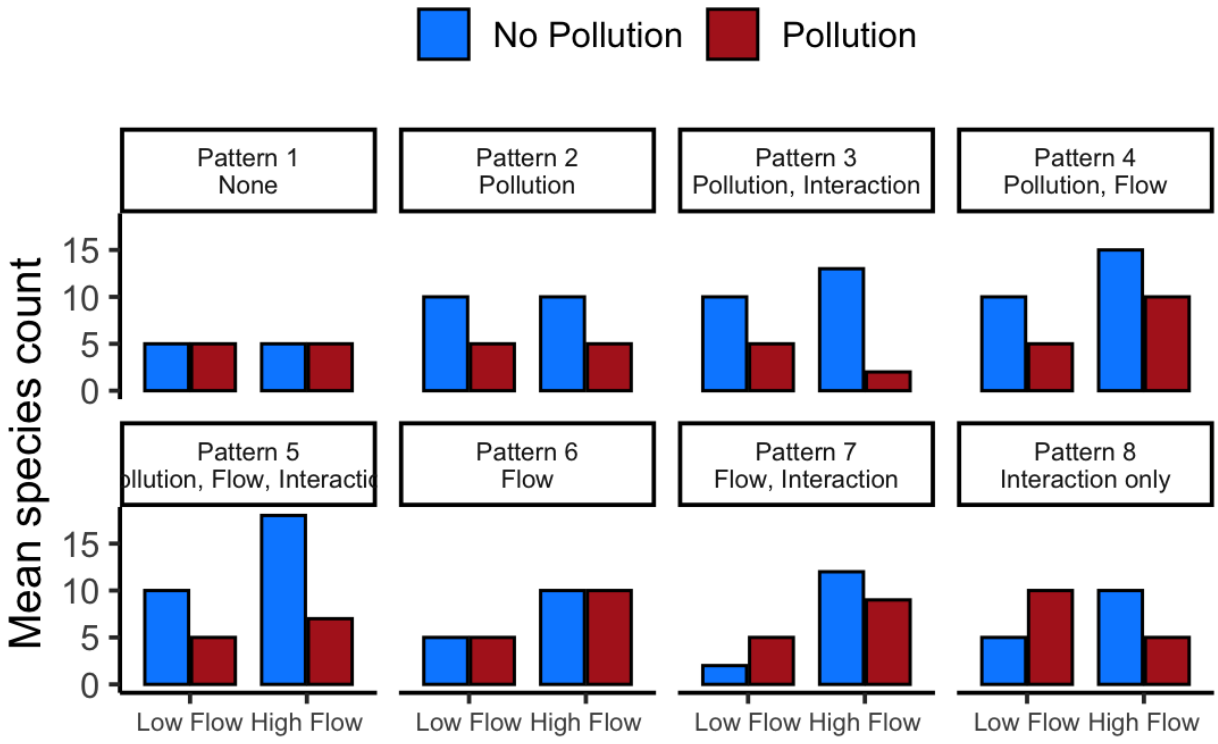


Figure 4: Eight example patterns of main effects and interactions in a pollution × flow design, numbered to match the pattern table above.

Figure 5 shows the same patterns in line graph form:

In the line graphs, interactions are easiest to see:

- **Parallel lines** mean no interaction (the pollution effect is the same at both flow rates).
- **Crossing lines** mean an interaction (the pollution effect differs by flow). The lines don't have to cross inside the plot: converging or diverging lines would cross if you extended them, so they count too. In a plot of sample cell means, the lines are almost never exactly parallel. The plot shows the pattern, and the interaction F test says whether the lines differ from parallel by more than chance would produce.

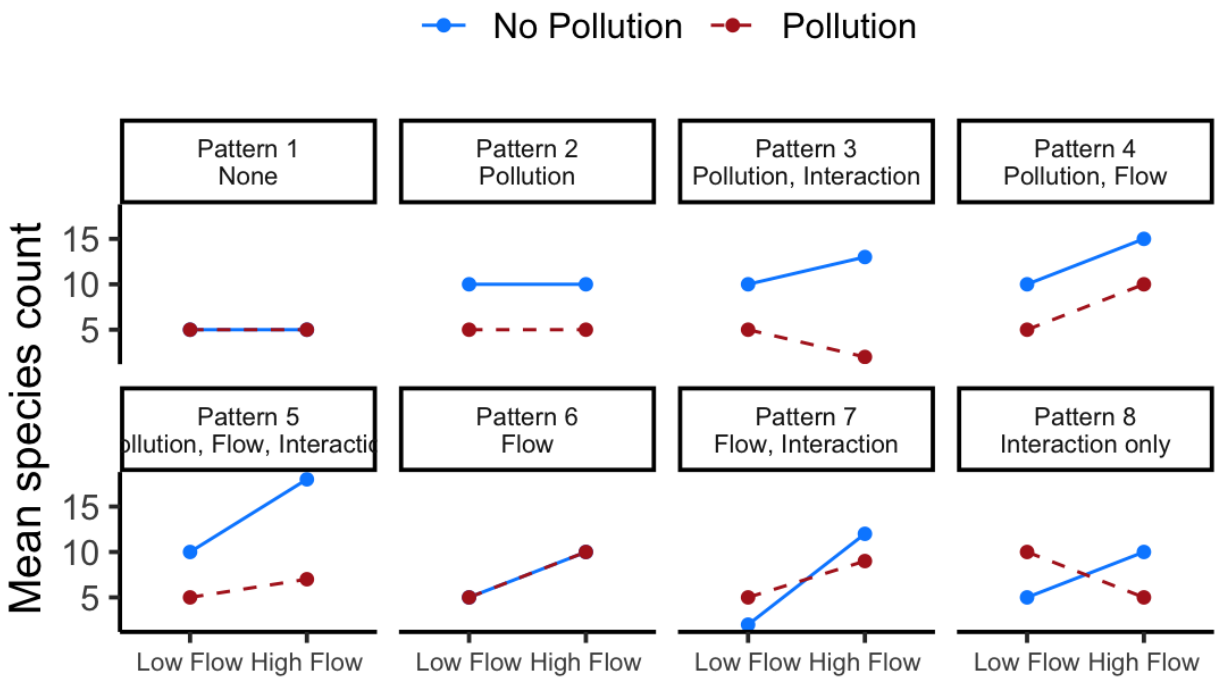


Figure 5: Line graphs showing eight possible outcomes for a pollution × flow design. In Patterns 1 and 6 the two lines lie on top of each other.

- The average gap between the two lines is the main effect of pollution, and the lines' average rise or fall from low to high flow is the main effect of flow.

> Tip: Which factor goes on the x-axis?
>
> The statistics are identical either way, since the same cell means get plotted regardless. Two rules of thumb make the picture easier to read. Put the factor with more levels on the x-axis, since one long line is easier to follow than several short, crowded ones. And draw the factor whose effect you're asking about as the lines, so its simple effect at each level of the other factor is the gap between the lines. That's why this chapter puts flow on the x-axis: the question is how the effect of pollution changes with flow.

Simple effects

When an interaction is present, the natural follow-up is to describe each simple effect, the effect of one factor at a single level of the other. Start with the Δ values from the cell means, as we did earlier:

$$\Delta_{\text{Pollution|Low}} = 9.6 - 3.6 = 6.0, \qquad \Delta_{\text{Pollution|High}} = 12.6 - 10.0 = 2.6$$

Those two numbers describe the interaction. Pollution costs 6.0 species under low flow and 2.6 under high flow. But saying "pollution reduced species count under low flow" is still a claim, and like any claim in this book, it needs its own test, not just a look at the plot.

Testing a simple effect asks whether pollution's effect holds at one specific level of flow, so it looks like a two-sample comparison: no pollution vs. pollution, at low flow only. The temptation is to pull out just the low-flow rows and run a fresh two-sample test on them, but that throws away the high-flow channels when estimating the error variance, so the estimate uses half the data and half the degrees of freedom. You already have a better estimate of that error variance, from all 20 channels: the SS_{Error} you computed for the omnibus ANOVA above.

The fix is to reuse that pooled estimate. MS_{Error} and its 16 degrees of freedom already describe the noise in this design better than any subset could, so the simple-effects test just plugs them into a two-sample t:

$$t = \frac{\bar{Y}_A - \bar{Y}_B}{\sqrt{MS_{\text{Error}}\left(\frac{1}{n_A} + \frac{1}{n_B}\right)}}, \qquad df = df_{\text{Error}}$$

Here A and B are the two pollution conditions being compared at one flow rate. Using $MS_{\text{Error}} = 1.60$ and $df = 16$ from the omnibus table computed earlier:

Table 10: Simple effects of pollution, tested separately within each flow rate, using the pooled error term from the full model.

Simple effect	Difference	t	df	p
Pollution, within Low Flow	6.0	7.50	16	< 0.001
Pollution, within High Flow	2.6	3.25	16	0.005

Both comparisons clear the critical value, $t_{crit}(16) = 2.12$ at $\alpha = 0.05$, two-tailed, with $t = 7.50$ under low flow and $t = 3.25$ under high flow. So the interaction here does not mean "pollution matters at low flow but

not at all at high flow." That pattern would show up as one significant row and one not. The interaction means "pollution matters at both flow rates, but far more at low flow."

When a factor has only two levels, as Pollution does here, the simple-effects test is the full answer. There is nothing left to compare.

Order of interpretation

When we conduct a two-way ANOVA, we always begin by reading the interaction term, because it decides what the main effects mean.

- If the interaction is *not significant*, examine the two main effects separately, interpreting each as an overall difference averaged across the other factor. If a significant main effect has three or more levels, follow it with Tukey HSD on that factor, exactly as in Chapter 8.
- If the interaction is significant, lead with the simple effects: how one factor behaves at each level of the other. "Pollution reduced species count sharply under low flow, and less under high flow" is a simple-effects sentence. "There was a main effect of pollution" on its own is not enough. Whether the main effect is still worth reporting depends on the direction of the simple effects:
  - If the simple effects all point the same way and differ only in size, the main effect still describes something real. Report it, qualified by the interaction: "There was a main effect of pollution, qualified by a Pollution × Flow interaction."
  - If one simple effect is near zero, or they point in opposite directions, the main effect is an average that describes no actual condition. Don't report it as a general effect of that factor.
- If *neither* the interaction nor either main effect is significant, there is nothing to follow up on. Do not run, or interpret, post-hoc tests on a result that never cleared the omnibus test in the first place.

The three examples below show each case of that rule.

A consistent main effect and an interaction

Figure 6 shows both a main effect and an interaction. Unpolluted channels are above polluted ones at both flow rates. That is the main effect. But the size of that gap is larger under low flow (12 vs. 5) than under high flow (14 vs. 11), so there's an interaction too.

There was a main effect of pollution, qualified by a Pollution × Flow interaction. The size of the pollution effect changed with flow rate, larger under low flow and smaller under high flow. The main effect is still real. Pollution is associated with lower species counts at both flow rates. Flow changes how much lower, not the direction.

An "averaged" main effect driven by an interaction

Figure 7 shows a starker pattern, with a large difference between polluted and unpolluted channels under low flow and no difference at all under high flow.

Yet an ANOVA on data with these cell means can still report a significant main effect of pollution when the error variance is small. Averaged across flow rates, unpolluted channels give $(10 + 5)/2 = 7.5$ species and polluted channels give 5, a main-effect difference of 2.5. That number is real, but it doesn't describe what pollution actually does. Pollution removes 5 species under low flow and none under high flow, not 2.5 everywhere. The "main effect" here is an artifact of averaging over a strong interaction.

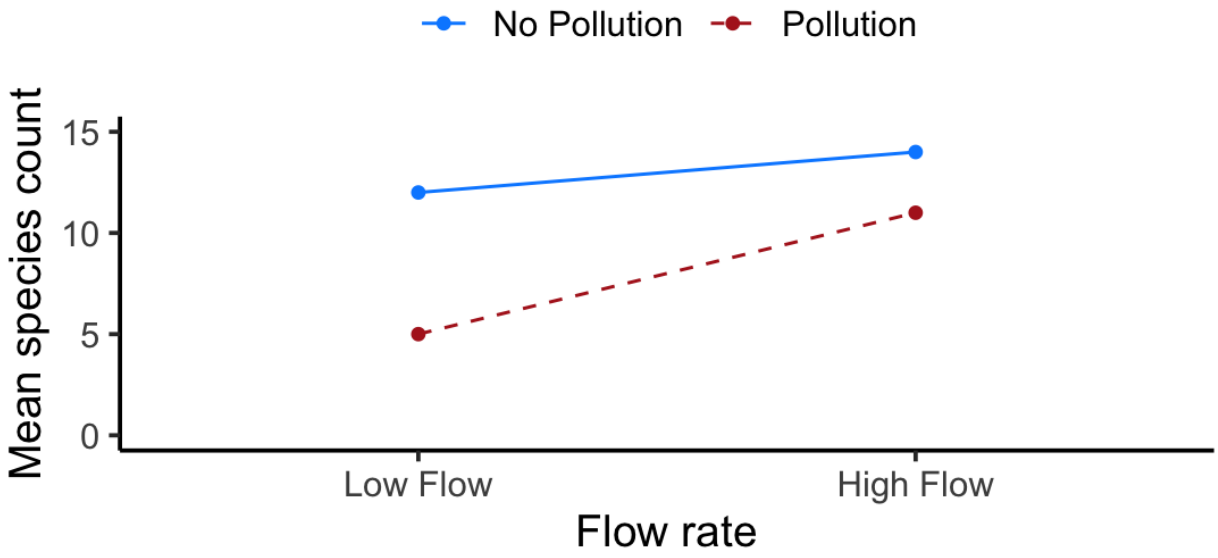


Figure 6: Pollution shows a general (main) effect, but the size of that effect differs by flow rate (interaction).

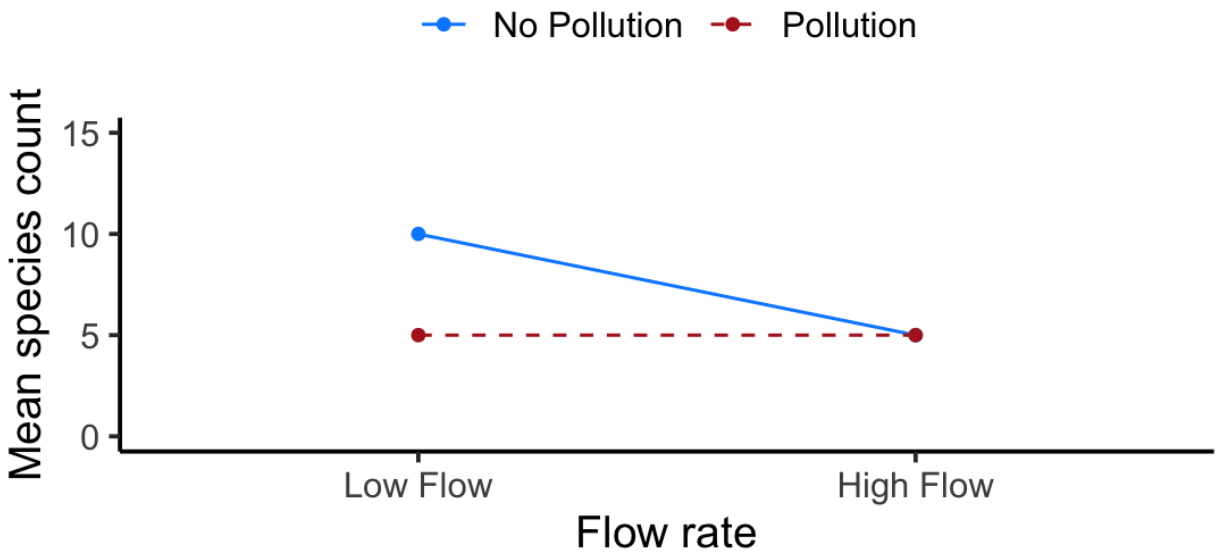


Figure 7: An apparent main effect of pollution arises from averaging over an interaction.

The ANOVA shows a significant main effect of pollution, but it should not be reported as a general effect, because the pollution effect was large under low flow and zero under high flow. In a case like this, it's fair to doubt whether pollution has a *general* effect on species count at all. It may only do damage under specific conditions, in combination with other factors, rather than on its own.

A main effect masked by an interaction

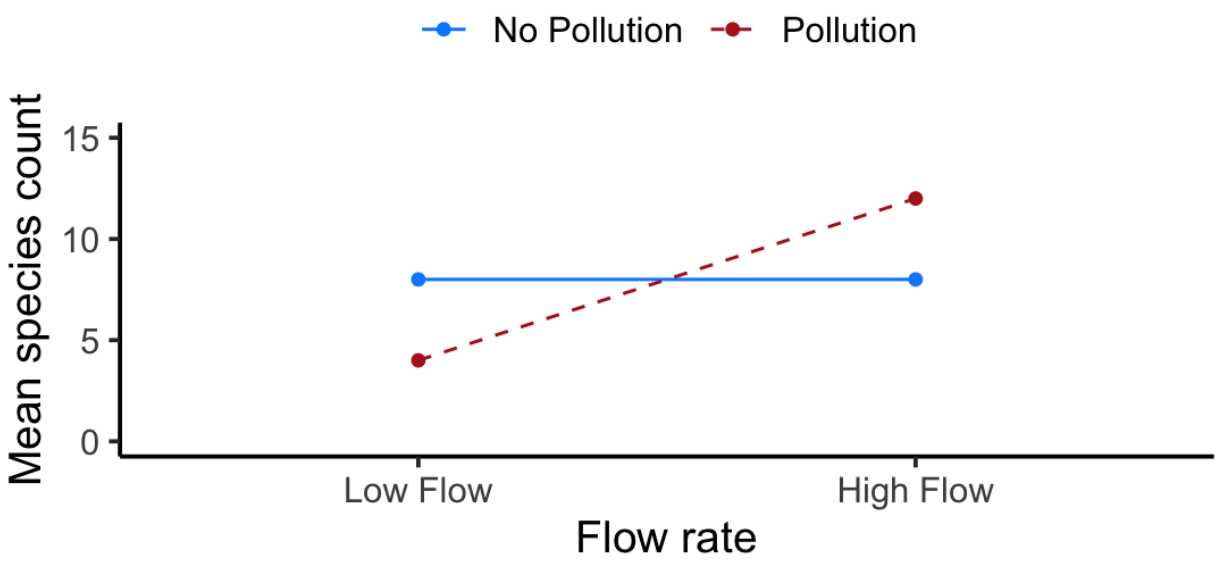


Figure 8: Pollution's effect reverses direction between flow rates, so the main effect averages to zero even though pollution matters at both.

Figure 8 shows the most misleading pattern of the three. Under low flow, pollution *reduces* species count (8 vs. 4). Under high flow, pollution *increases* it (8 vs. 12). Both are real, sizable effects, just pointed in opposite directions.

Average across flow rates to get a main effect, and they cancel: unpolluted channels give $(8 + 8)/2 = 8$, polluted channels give $(4 + 12)/2 = 8$, a main-effect difference of exactly 0. An ANOVA on data with these cell means would report no main effect of pollution at all, because its two opposite effects average each other out.

This is the case the "read the interaction first" rule exists for. A non-significant main effect here would be true and misleading. Pollution has no effect *on average* and a large effect *in each condition*. Reporting only the main effect would suggest pollution is irrelevant, when it has a large effect at both flow rates, in opposite directions.

Table 11 puts the three cases and the stream experiment side by side.

Table 11: The same rule applied to four sets of cell means.

Case	Simple effect of pollution (low flow, high flow)	Main effect of pollution	What to report
Consistent main effect and an interaction (Figure 6)	7, 3	5	Simple effects, plus the main effect with its qualification
Averaged main effect (Figure 7)	5, 0	2.5, which matches neither condition	Simple effects only
Masked main effect (Figure 8)	4, −4	0, which hides two real effects	Simple effects only
Stream experiment (Figure 9)	6.0, 2.6	4.3	Simple effects, plus the main effect with its qualification

Reporting the stream result

Back to the stream experiment. Figure 9 plots its four cell means. Pollution lowers species count at both flow rates, so the lines never cross. They converge, because the gap under high flow is smaller. That is the first case of the rule, so we report the simple effects and the main effect of pollution with its qualification attached.

! Results reporting: factorial ANOVA

Lead with the interaction, then give the simple effects. Report a main effect only when the simple effects point the same way, and then with its qualification attached. Every effect gets its own effect size, labeled so the reader knows whether it is partialed.

> The effect of pollution on macroinvertebrate species count depended on flow rate, $F(1, 16) = 9.03$, $p = 0.008$, $\eta_p^2 = 0.36$. Under low flow, polluted channels supported 6.0 fewer species than unpolluted ones (3.6 vs. 9.6), $t(16) = 7.50$, $p < 0.001$. Under high flow that gap narrowed to 2.6 species (10.0 vs. 12.6), $t(16) = 3.25$, $p = 0.005$. Both main effects were also significant, pollution $F(1, 16) = 57.78$, $p < 0.001$, $\eta_p^2 = 0.78$, and flow $F(1, 16) = 69.03$, $p < 0.001$, $\eta_p^2 = 0.81$, but each should be read in light of the interaction:

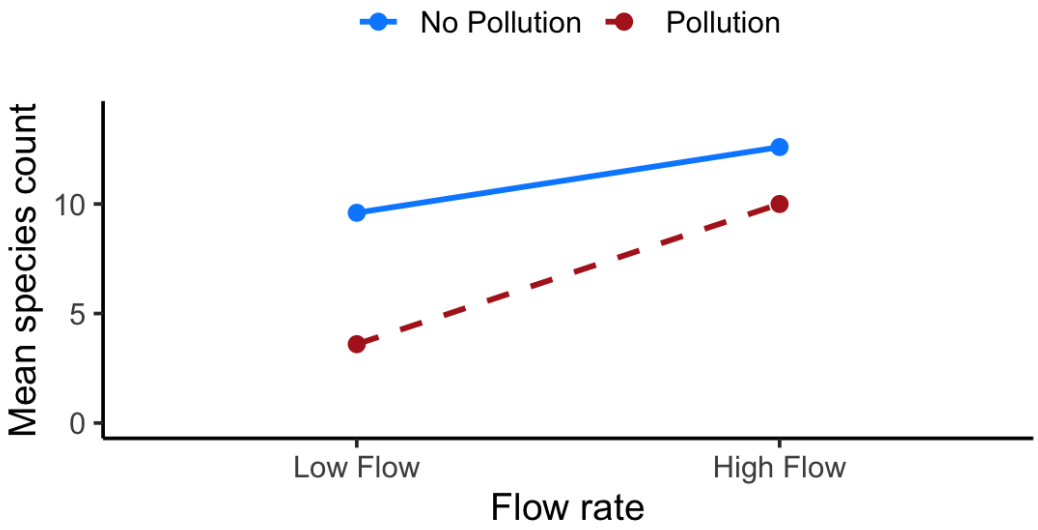


Figure 9: Cell means from the stream experiment. Pollution lowers species count at both flow rates, by more under low flow.

> pollution costs a stream far more species when the water is moving slowly.

Table 12 collects the two-factor ANOVA on one card.

Table 12: Test card: two-factor ANOVA.

Purpose	Do two factors, crossed in every combination, each change the mean, and does the effect of one depend on the level of the other?
Hypotheses	No main effect of A No main effect of B No interaction, $H_0 : (\mu_{A1,B1} - \mu_{A2,B1}) = (\mu_{A1,B2} - \mu_{A2,B2})$
Assumptions	Similar spread in every cell and approximately normal residuals
Test statistic	$F = MS_{\text{Effect}}/MS_{\text{Error}}$ for each effect. df: $a-1$ (A), $b-1$ (B), $(a-1)(b-1)$ (interaction), $N-ab$ (error)
Report	Interaction first, $F(df, df_{\text{Error}})$, p, η_p^2 (labeled as partial). If it is significant, the simple effects, and a main effect only when the simple effects point the same way. If it is not, each main effect, with Tukey HSD for a significant factor with three or more levels

More complicated designs

Up until now we have focused on the simplest case for factorial designs, the 2x2 design, with two IVs, each with 2 levels. The same logic extends to larger designs.

3x2 design

This design works like the 2x2, with one more level added to one factor. The test structure is the same (two main effects and one interaction). The interaction is harder to interpret.

The same stream example extends to a 3x2. Suppose we examine macroinvertebrate species count (DV) under three flow rates (low, medium, high) crossed with two pollution conditions (no pollution vs. pollution).

- The main effect of pollution tests whether, on average, polluted and unpolluted channels differ in species count across all flow rates.
- The main effect of flow tests whether mean species count differs among flow rates.
- The interaction tests whether the pollution effect (unpolluted − polluted) changes with flow rate.

In a 2x2 design, you can easily spot interactions by looking for non-parallel lines. In a 3x2, the idea is the same, but with three points per line the patterns are harder to read.

You have three hypotheses in a 3x2 factorial design:

Hypothesis 1: Main effect of flow

- **Null ($H_{0,\textbf{Flow}}$):** $\mu_{\text{Low}} = \mu_{\text{Medium}} = \mu_{\text{High}}$
- **Alternative ($H_{a,\textbf{Flow}}$):** At least one flow rate mean differs from the others.

Hypothesis 2: Main effect of pollution

- **Null ($H_{0,\textbf{Pollution}}$):** $\mu_{\text{NoPoll}} = \mu_{\text{Poll}}$
- **Alternative ($H_{a,\textbf{Pollution}}$):** Mean species count differs between polluted and unpolluted channels.

Hypothesis 3: Pollution × flow interaction

- **Null ($H_{0,\textbf{int}}$):** The pollution effect is the same across all flow rates. $[(\mu_{\text{NoPoll,Low}} - \mu_{\text{Poll,Low}}) = (\mu_{\text{NoPoll,Medium}} - \mu_{\text{Poll,Medium}}) = (\mu_{\text{NoPoll,High}} - \mu_{\text{Poll,High}})]$
- **Alternative ($H_{a,\textbf{int}}$):** The pollution effect differs across flow rates (i.e., at least one difference between differences is not equal).

One factorial ANOVA is better than running a separate t-test of polluted vs. unpolluted channels at each of the three flow rates, because you have more statistical power (higher degrees of freedom) for the tests of interest, and you get a formal test of the interaction between factors which is often scientifically interesting.

Figure 10 shows example means for this design.

The figure shows invented means for all six conditions.

1. **Main effect of pollution.** Averaged across flow rates, unpolluted channels support higher species counts (10 vs. 7.3). The "No Pollution" line sits above the "Pollution" line at low and medium flow, and the two coincide at high flow.

2. **Main effect of flow.** Averaged over pollution condition, species counts are lowest at low flow (8) and higher at medium and high flow (9 each).

3. **Interaction between pollution and flow.** An interaction occurs when the effect of one IV depends on the levels of another. The pollution gap is 6 species at low flow, 2 at medium, and 0 at high. So the size of the pollution effect changes with flow rate, which is also why the pollution main effect above describes no single flow condition well. You would have to conduct an inferential test on the interaction term to see if these differences were likely or unlikely to be due to sampling error.

If there were no interaction and no main effect of flow, we would see something like the pattern in Figure 11.

Figure 12 shows a different interaction pattern.

The means show an interaction here too. The pollution effect is measured three times, once at each flow rate. It is the same at low and medium flow but much smaller at high flow, so its size depends on the level of flow.

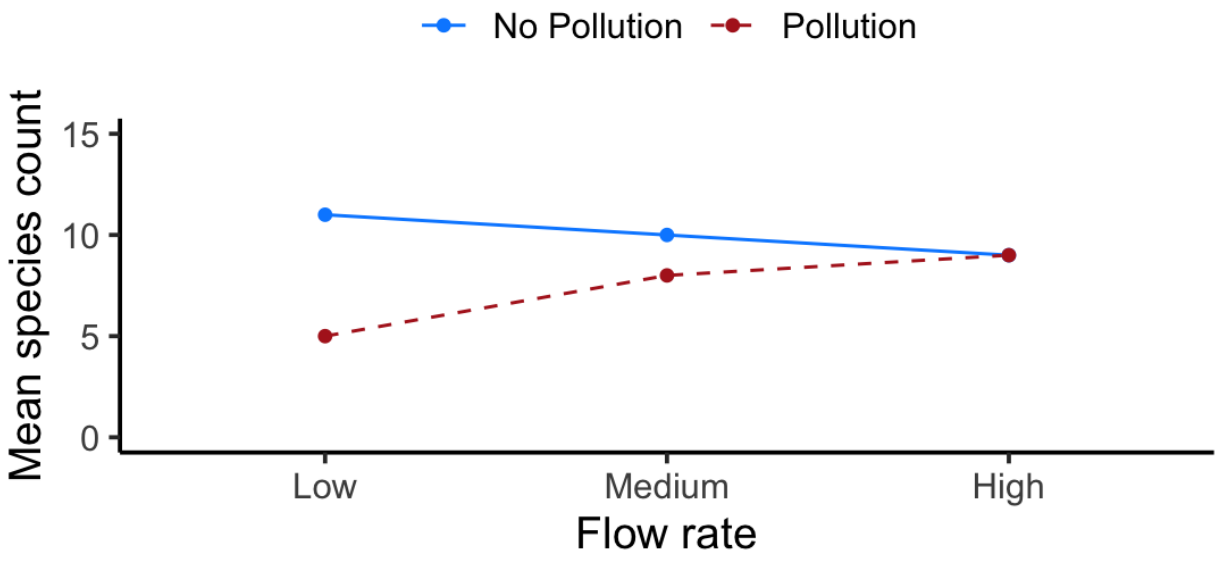


Figure 10: Example means for a 3 × 2 pollution × flow design.

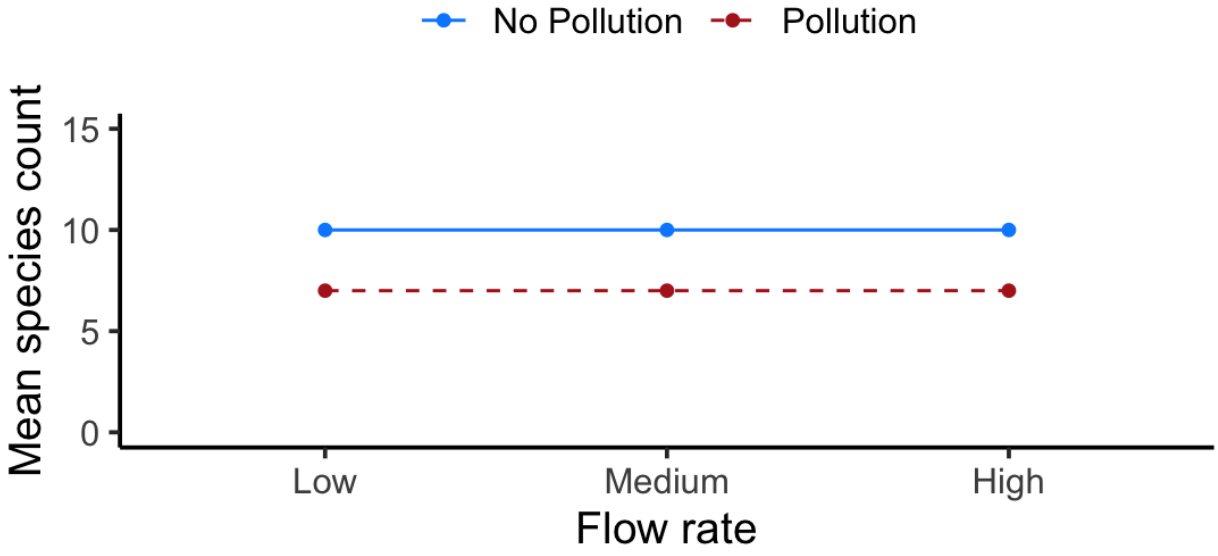


Figure 11: Example means for a 3 × 2 design with only one main effect.

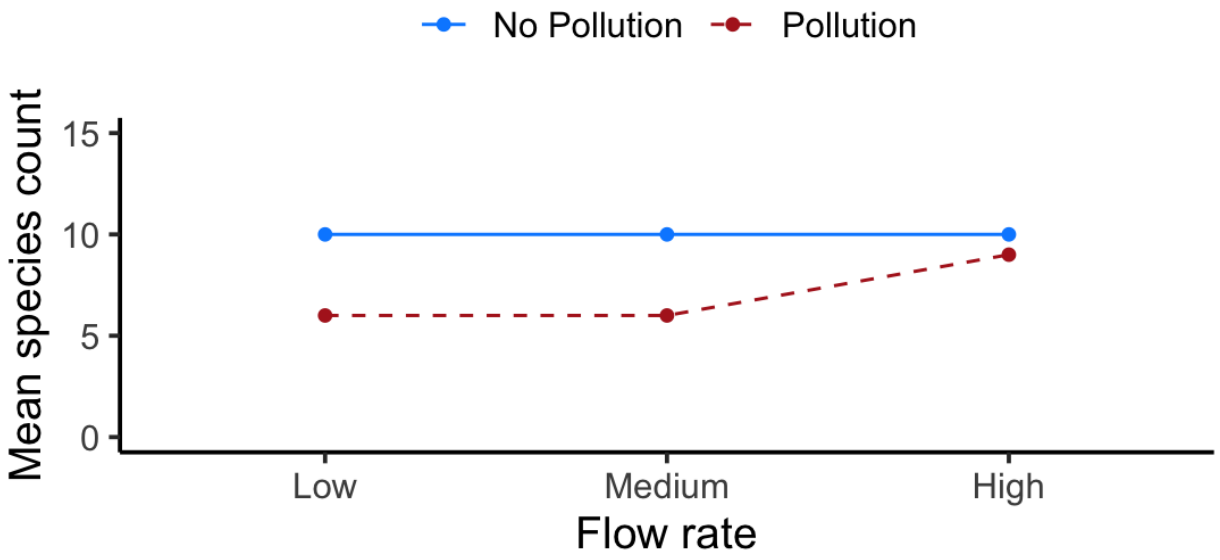


Figure 12: Example means for a 3 × 2 design showing a different interaction pattern.

Designs beyond this course

Larger factorial designs use the same logic as the 2x2. But complexity grows quickly, both in the number of effects and in how hard they are to interpret.

2x2x2 (three-factor) designs

A 2x2x2 design includes three independent variables, each with two levels. That creates:

- three main effects (one per variable)
- three **two-way interactions** (A×B, A×C, B×C), each an interaction between two of the three variables
- one three-way interaction (A×B×C)

You can think of it as two 2x2 plots side by side, one for each level of the third variable.

A **three-way interaction** means that one of the two-way interactions changes across the third variable. For example, the pollution × flow interaction might look different in summer than in winter.

Interactions can themselves interact. The math is identical in structure, but the number of effects grows fast. A 2x2 tests 3 effects, and a 2x2x2 tests 7. The graphs also become difficult to interpret. In applied research, most studies stop at two factors unless there's a strong reason to add more.

Mixed designs

A **mixed design** has at least one between-subjects factor and at least one within-subjects factor. For example, sites or treatments could be a between-subjects factor, and repeated measures over time a within-subjects factor.

The interpretation of main effects and interactions is the same, but the error terms used in the F-tests differ, and the data structure must account for repeated measures. These models are important in longitudinal and environmental time-series studies, but they're beyond our current scope.

Factorial ANOVA can handle many experimental structures, but each added factor makes the results harder to interpret, so use only the factors your question needs.

Chapter Summary

Why it matters. Single-factor designs can tell you whether something has an effect. Factorial designs tell you what *changes* that effect, which is usually the more interesting scientific question, and often the more realistic one for environmental systems where drivers act together. An interaction changes what every other term in the table means, so interpreting the table is part of the analysis. Chapter 11 replaces one factor with a continuous covariate.

Core ideas

- **Fully crossed means every combination.** A factorial design pairs every level of each IV with every level of the others. The notation says the structure: 2x2 is four conditions, 3x2 is six.
- **A 2x2 tests three things.** One main effect per factor, plus their interaction. That means three null hypotheses, and you should be able to state all three in plain language before running anything.
- **An interaction is a difference between differences.** It asks whether the effect of one factor changes across levels of another. The interaction sum of squares computes this difference between differences.
- **The partition extends the one-way logic.** SS_{Total} now splits into each main effect, the interaction, and error (SS_{Error}, Chapter 8's SS_{Within} in the one-way ANOVA). Each effect gets its own F, all sharing MS_{Error} as the denominator. Working one table through by hand shows that the interaction term is the only new piece.
- **Three tests, three effect sizes.** η^2 divides each effect's SS by SS_{Total}, and partial η_p^2 divides it by that effect plus error. Partial values are at least as large and do not sum to 1, so label which one you are reporting. A significant effect can still be a small one, which is exactly the case for this chapter's interaction.
- **Read the interaction first.** If it's significant, lead with simple effects: what factor A does at each level of B, as Δ values from the cell means, tested with the pooled MS_{Error}. A main effect can still be reported, qualified, when the simple effects all point the same way. When one is near zero or they reverse, the main effect describes no actual condition. When the interaction is not significant, the main effects stand on their own, with Tukey HSD for a factor with three or more levels.
- **Learn to read the picture.** In a line plot, parallel lines mean no interaction, and lines that cross, inside the plot or if extended, mean an interaction. The average gap between the lines is one main effect, and the lines' average rise or fall is the other. There are eight possible patterns in a 2x2, and recognizing them by eye is faster than parsing the table.
- **An averaged main effect can be an artifact.** A factor that matters a lot in one condition and not at all in another can still produce a "significant main effect" once you average across conditions. That number is correct arithmetic, but it matches the pollution effect at neither flow rate.
- **Bigger designs, same tests, harder reading.** A 3x2 has the same three tests as a 2x2, but three points per line are harder to read than two. Beyond that, a three-factor design brings three main effects, three two-way interactions, and a three-way interaction, and interpretability falls off fast. Use only as many factors as your question actually needs.